

Walmart

INDEPENDENT CANDIDATE VISION FOR WALMART

Earn the right to standardize.

A tactical-empathy-first entry plan for Global Governance Digital Strategy.

Days 1-30 - Listen and map

Build a shared, evidence-based view of decisions, work, authority, and friction before proposing a new governance mechanism.

PRIMARY OUTCOME

What changes by day 30

Leadership has one map of the current portfolio and the decisions that govern it.

- Meet CLO staff, Governance leaders, Technology, Finance, Architecture, Privacy, Security, and delivery teams.
- Inventory active and proposed AI use cases, spend, dependencies, vendors, controls, and accountable owners.
- Map intake, prioritization, approval, exception, escalation, release, monitoring, and retirement decisions.
- Shadow representative workflows across Legal, Compliance, Ethics, Investigations, and Security.
- Separate prudent review time from avoidable queue time and rework.
- Baseline decision latency, evidence completeness, adoption burden, and recurring bottlenecks.

QUESTIONS TO ANSWER

- Where does AI demand enter, and how does it become portfolio work?
- Which decisions are enterprise-wide, and which require local legal judgment?
- Where are ownership, funding, or release authority ambiguous?
- Which controls are repeatedly rebuilt rather than reused?
- Where does adoption fail after technical approval?
- Which existing mechanisms already work and should be protected?

STAKEHOLDER LISTENING MAP

Stakeholder	What to learn
CLO and Governance leadership	Strategic priorities, risk appetite, decision pain, and executive visibility needs
Legal, Compliance, Ethics, Investigations	Substantive authority, workflow friction, exceptions, and evidence standards
Privacy and Security	Data boundaries, access, logging, monitoring, and escalation expectations
Technology and Architecture	Platforms, dependencies, reusable services, and delivery constraints
Finance and business owners	Investment logic, baselines, value evidence, and operating ownership

DISCOVERY ARTIFACTS

- Decision-rights map
- Stakeholder topology
- Portfolio and spend inventory
- Use-case risk segmentation
- Workflow friction and rework map
- Control reuse map
- Adoption burden map
- Baseline scorecard

EARLY TRUST SIGNALS

Clear listening notes, visible follow-through, transparent assumptions, disciplined meeting load, and removal of one avoidable operating friction without weakening a necessary control.

Days 31-60 - Make decisions visible

Co-design the minimum portfolio contract and decision architecture with the people who own the consequences.

PRIMARY OUTCOME

What changes by day 60

Material AI work is visible in one portfolio view with explicit ownership and a defined next decision.

- Co-design the seven-clause Scannable AI Contract with governance, technology, finance, and operating leaders.
- Define risk-sensitive stage gates rather than forcing every use case through the same review burden.
- Create a portfolio view showing sponsor, owner, value thesis, risk posture, dependencies, evidence, and next decision date.
- Clarify who may Fund, Shape, Hold, Stop, release, escalate, or grant an exception.
- Align Finance on investment, capacity, and value-realization logic.
- Select two contrasting use cases for controlled proof.

MINIMUM DECISION CONTRACT

Clause	Question made explicit
Outcome	Which workflow, decision, risk, or obligation improves?
Legal basis	What authority, obligation, policy, or permitted use supports the work?
Risk and harm	What can go wrong, who is affected, and what escalates?
Data and security	What data is necessary, protected, retained, logged, or prohibited?
Human authority	Which decisions remain human and where may automation act?
Operating ownership	Who owns performance, exceptions, adoption, maintenance, and retirement?
Evidence and economics	What proof earns the next investment or triggers a stop?

DECISION RIGHTS TO CLARIFY

Decision	Required clarity
Intake	Who qualifies demand and what minimum context is required?
Prioritization	How value, risk, reuse, urgency, and capacity are balanced
Release	Who accepts residual risk and operational readiness
Exception	Who may deviate, for how long, and with what evidence
Funding	What evidence earns additional investment
Retirement	Who stops low-value, obsolete, or unmanaged work

CONTROLLED-PROOF PAIR

Use case A - lower-risk, high-adoption potential

Tests whether the contract accelerates a useful workflow without adding unnecessary governance load.

Use case B - sensitive data or authority boundary

Tests whether the mechanism exposes ambiguity early enough to reshape or hold the work responsibly.

EXIT CONDITION

The mechanism is ready for proof when each use case has a named sponsor, accountable operating owner, explicit human decision boundary, evidence plan, and next decision date.

Days 61-90 - Run controlled proof

Test whether the operating mechanism improves real decisions - not whether it produces more documentation.

PRIMARY OUTCOME

What changes by day 90

Two contrasting use cases produce evidence about decision quality, control sufficiency, adoption burden, and value.

- Apply the contract to two use cases with different risk, data, and authority profiles.
- Test evaluation, logging, monitoring, escalation, fallback, and human-review mechanisms.
- Measure decision latency, rework, evidence completeness, and adoption burden.
- Run an executive portfolio review that results in a real Fund, Shape, Hold, or Stop decision.
- Capture reusable controls, decision templates, architecture patterns, and unresolved exceptions.
- Document what the mechanism accelerated, what it protected, and where it created unnecessary burden.

EVIDENCE PLAN

Category	Evidence sought
Decision latency	Time from intake to a named posture, segmented by risk and complexity
Control evidence	Quality of evaluation, logging, monitoring, escalation, and human review
Adoption burden	Training, workarounds, exceptions, role changes, and sustained use
Value evidence	Observed outcome against an agreed baseline and operating cost
Exception quality	Whether exceptions reveal policy, architecture, ownership, or data gaps

TEST SEQUENCE

Before release

Confirm the decision contract, evidence threshold, human authority, fallback, owner, and monitoring plan.

During controlled use

Observe real workflow behavior, exceptions, data handling, review load, and user workarounds.

After evidence review

Change funding, design, control, release, ownership, or retirement based on what was learned.

EXECUTIVE PORTFOLIO REVIEW OUTPUT

- Named posture and rationale
- Accountable owner
- Next evidence requirement
- Human decision boundary
- Control or architecture dependency
- Funding or capacity implication
- Next decision date

FAILURE MODES TO SURFACE

- Approval without operating ownership
- Strong demo with weak adoption path
- Control burden disproportionate to risk
- Automation crossing an unclear human-authority boundary
- Value claim without a defensible baseline

EXIT CONDITION

Evidence has changed at least one material portfolio decision and the review has identified which elements should be reused, revised, localized, or stopped.

Days 91-120 - Scale the mechanism

Turn proven decisions, controls, and operating practices into a two-quarter roadmap without flattening legitimate local judgment.

PRIMARY OUTCOME

What changes by day 120

Leadership has a prioritized roadmap tied to value, trust, talent, capacity, and reusable governance capabilities.

- Publish reusable decision, control, evaluation, and ownership patterns.
- Align capacity and talent to the prioritized portfolio roadmap.
- Define executive cadence, escalation paths, and portfolio health indicators.
- Set value-realization and adoption reviews with Finance and business owners.
- Establish a learning backlog: what to standardize, localize, automate, or stop.
- Translate unresolved exceptions into policy, platform, data, talent, or operating-model work.

TWO-QUARTER ROADMAP WORKSTREAMS

Workstream	Roadmap intent
Portfolio governance	Single intake, prioritization, posture, ownership, and decision cadence
Reusable controls	Shared evaluation, logging, monitoring, escalation, and evidence patterns
Platform and data	Architecture and data capabilities that reduce duplicated review and delivery work
Adoption and talent	Role clarity, training, change burden, support, and capacity alignment
Measurement and learning	Value realization, decision latency, evidence quality, reuse, and retirement

EXECUTIVE CADENCE

Cadence	Purpose
Monthly portfolio review	Fund, Shape, Hold, Stop, remove dependencies, and reallocate capacity
Quarterly value review	Compare realized outcomes with investment, adoption, and operating cost
Exception learning review	Convert recurring exceptions into reusable policy, control, or platform work
Talent and capacity review	Match specialist expertise and delivery capacity to the highest-value portfolio needs

120-DAY DELIVERABLES

- Governance portfolio view
- Scannable AI Contract
- Risk-sensitive stage gates and decision-rights map
- Executive scorecard and review cadence
- Reusable control and evaluation library
- Two-quarter roadmap
- Learning backlog and retirement candidates

SUCCESS AT DAY 120

Leadership can see the portfolio, understand the next decision for every material initiative, distinguish value from activity, preserve substantive governance authority, and scale expertise through a small number of trusted operating mechanisms.

CONTINUING QUESTIONS

- Did evidence change funding, design, release, ownership, or retirement?
- Which controls became reusable enterprise capabilities?
- Where does global consistency still conflict with local judgment?
- Which work should stop to release capacity?